

What is a Database System?

- A collection of:
 - Software
 - Files
 - Abstractions
 - Views
- Enables users to
 - Access data
 - Manipulate data
 - Otherwise manage data
 - Create reports

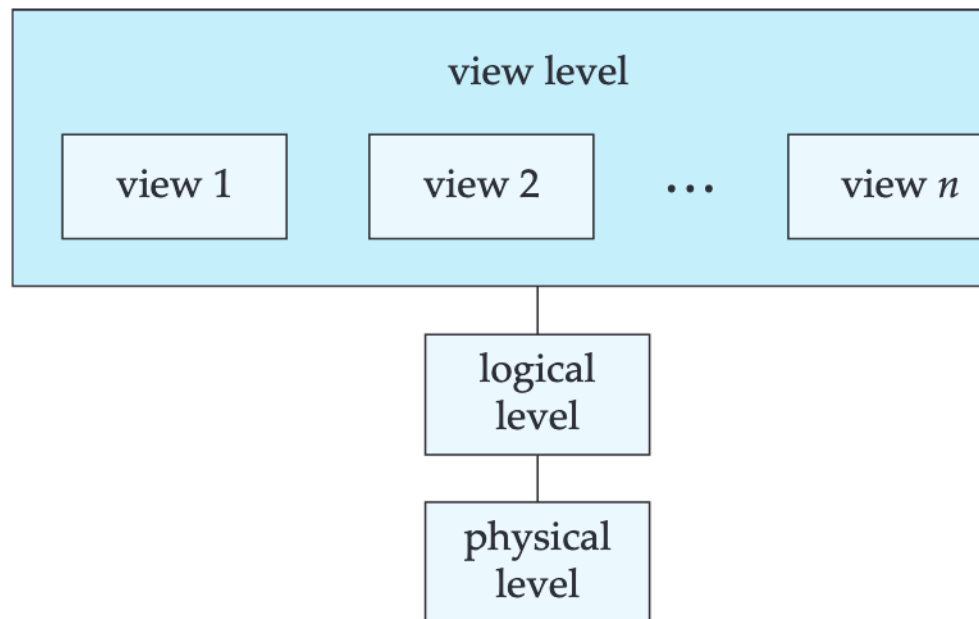


Why is Abstraction Important?

A major purpose of a database system is to provide users with an abstract view of the data.

- Data models
 - A collection of conceptual tools for describing data, data relationships, data semantics, and consistency constraints.
- Data abstraction
 - Hide the complexity of data structures to represent the database from users through several levels of abstraction.

An Architecture For a Databas

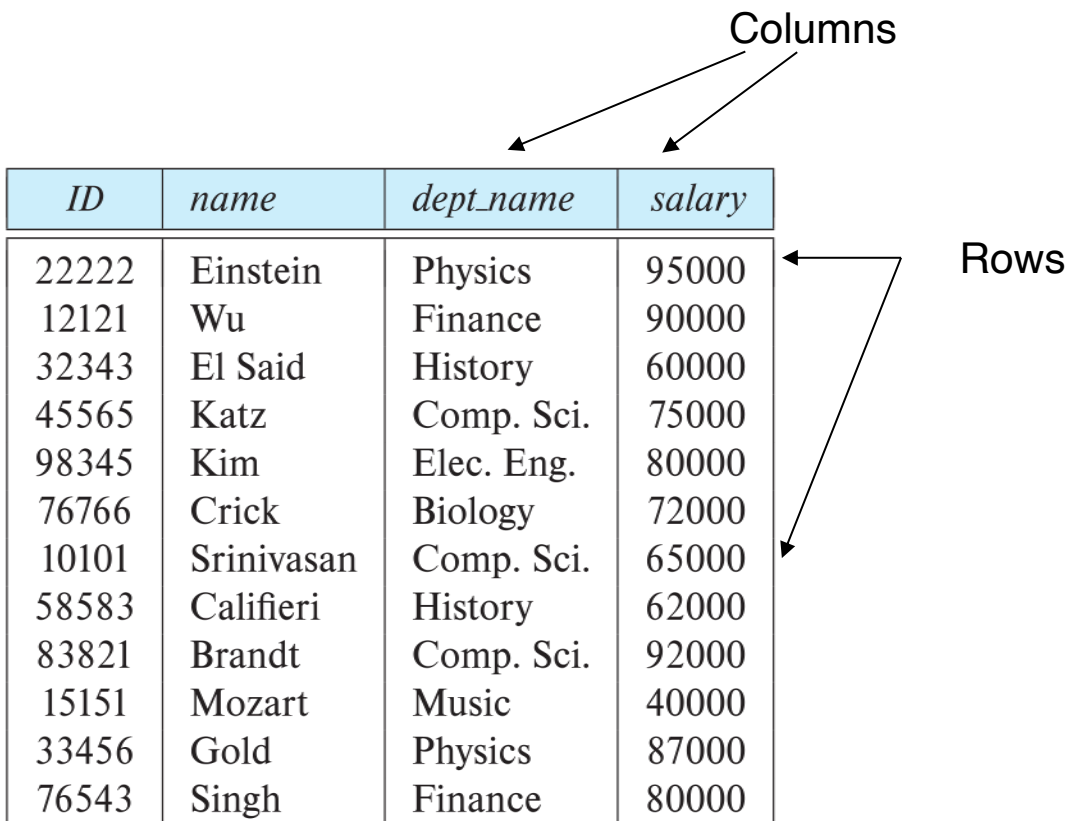


Data Models

- A collection of tools for describing
 - Data
 - Data relationships
 - Data semantics
 - Data constraints
- Relational model (data is presented as tables)
- Entity-Relationship data model (mainly for databases)
 - Object-based data models (Object-oriented and Object-relational)
- Semi-structured data model (XML, JSON)
- Other older models:
 - Network model
 - Hierarchical model

Relational Model

- All the data is stored in various tables.
- Example of tabular data in the relational model



The diagram illustrates a table structure. Two arrows labeled "Columns" point to the header row of the table. Two arrows labeled "Rows" point to the data rows of the table.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

A Sample Relational Database

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

<i>dept_name</i>	<i>building</i>	<i>budget</i>
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

Instances and Schemas

- Similar to types and variables in programming languages
- **Logical Schema** – the overall logical structure of the database
 - Example: The database consists of information about customers and accounts in a bank and the relationships between them
 - Analogous to type information of a variable
- **Physical schema** – the overall physical structure of the database
- **Instance** – the actual content of the database at a point in time
 - Analogous to the value of a variable

Physical Data Independence

- **Physical Data Independence** – the ability to modify physical schema without changing the logical schema
 - Applications depend on the logical schema
 - In general, the interfaces between the various components should be well defined so that changes in some parts do not seriously influence others.

So how do we set up and access the data
in a relational database?

Data Definition Language (DDL)

- Specification notation for defining the database schema

Example: **create table** *instructor* (
 ID **char**(5),
 name **varchar**(20),
 dept_name **varchar**(20),
 salary **numeric**(8,2))

- DDL compiler generates a set of table templates stored in the ***data dictionary***
- Data dictionary contains metadata (i.e., data about data)
 - Database schema
 - Integrity constraints
 - Primary key (ID uniquely identifies instructor)
 - Authorization
 - Who can access what

Data Manipulation Language (DML)

- Language for accessing and updating the data organized in an appropriate data model
 - DML also known as query language
- There are basically two types of data-manipulation languages
 - **Procedural DML** -- require a user to specify what data are required and how to get those data.
 - **Declarative DML** -- require a user to specify what data are required without specifying how to get those data.
- Declarative DMLs are usually easier to learn and use than procedural DMLs.
- Declarative DMLs are also referred to as non-procedural.
- The portion of a DML that involves information retrieval is called **query** language.

SQL Query Language

- SQL query language is nonprocedural. A query takes tables as input and always returns a single table.
- Example to find all instructors in Comp. Sci. dept

```
select name
from instructor
where dept_name = 'Comp. Sci.'
```
- To be able to compute complex functions SQL is usually used in some higher-level language
- Application programs generally access databases through
 - Language extensions to allow embedded SQL
 - Application program interface (e.g., ODBC/JDBC) to send SQL queries to a database

Database Access from Application P

- SQL does not support actions such as input from u to displays, or communication over the network.
- Such computations and actions must be written in **language**, such as C/C++, Java or Python, with embedded SQL queries that access the data in the database.
- **Application programs** -- are programs that are used to interact with the database in this fashion.

Designing a Relational Database

Database Design

The process of designing the general structure of the database

- Logical Design – Deciding on the database schema.
 - Business decision – What attributes should we need in the database?
 - Computer Science decision – What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
- Physical Design – Deciding on the physical layout of the database

The Database Engine

Database Engine

- A database system is partitioned into modules that each of the responsibilities of the overall system.
- The functional components of a database system are divided into
 - The storage manager,
 - The query processor component,
 - The transaction management component.

Storage Manager

- A program module that provides the interface between high level data stored in the database and the applications and queries submitted to the system.
- The storage manager is responsible to the following:
 - Interaction with the OS file manager
 - Efficient storing, retrieving and updating of data
- The storage manager components include:
 - Authorization and integrity manager
 - Transaction manager
 - File manager
 - Buffer manager

Storage Manager (Cont.)

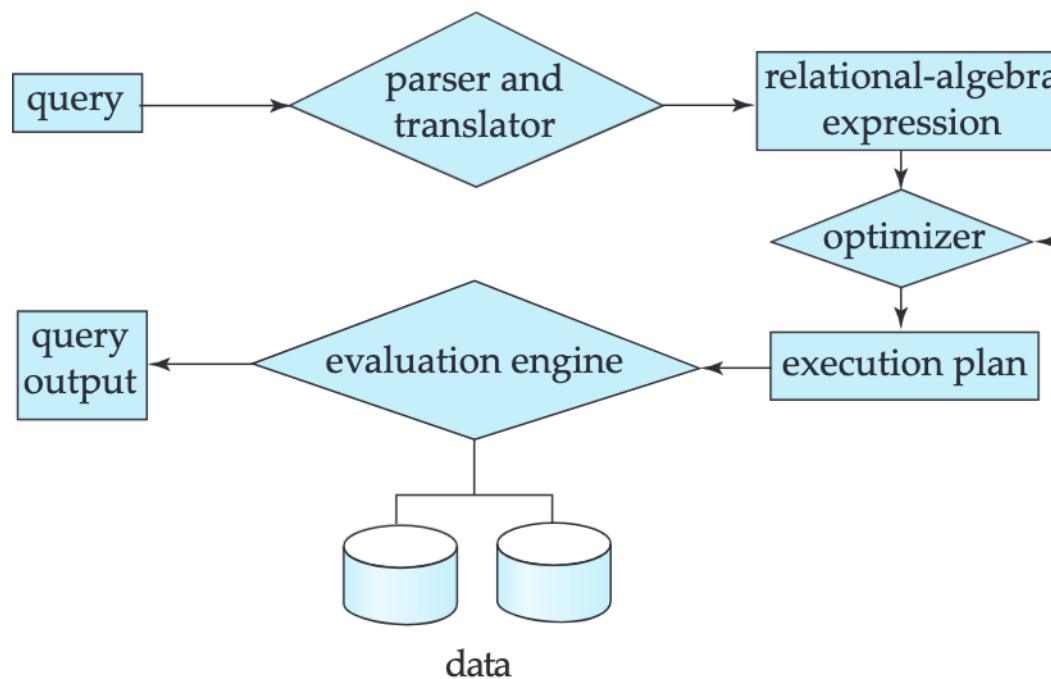
- The storage manager implements several data structures as part of the physical system implementation:
 - Data files -- store the database itself
 - Data dictionary -- stores metadata about the structure of the database, in particular the schema of the database
 - Indices -- can provide fast access to data items. A database index provides pointers to those data items that hold a particular value.

Query Processor

- The query processor components include:
 - DDL interpreter -- interprets DDL statements and the definitions in the data dictionary.
 - DML compiler -- translates DML statements in a language into an evaluation plan consisting of low-level instructions that the query evaluation engine understands.
 - The DML compiler performs query optimization; it picks the lowest cost evaluation plan from among various alternatives.
 - Query evaluation engine -- executes low-level instructions generated by the DML compiler.

Query Processing

1. Parsing and translation
2. Optimization
3. Evaluation



Transaction Management

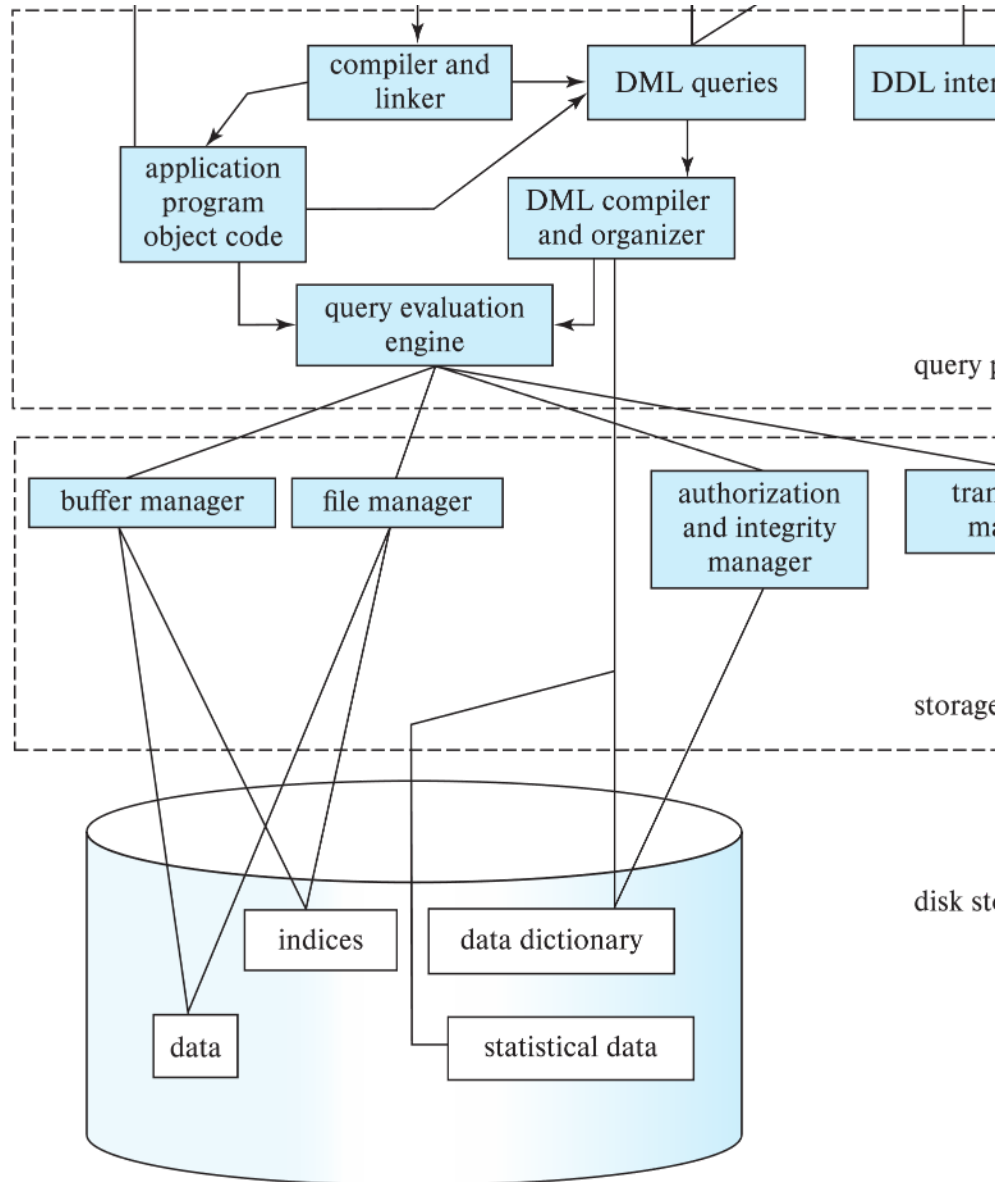
- A **transaction** is a collection of operations that perform a single logical function in a database application
- **Transaction-management component** ensures that the database remains in a **consistent** (correct) state despite system failures (e.g., power failures and operating system crashes) and transaction failures.
- **Concurrency-control manager** controls the interaction among the concurrent transactions (**isolation**), to ensure the consistency of the database.

A	Atomic
C	Consistent
I	Isolated
D	Durable

Database Architecture

- Centralized databases
 - One to a few cores, shared memory
- Client-server,
 - One server machine executes work on behalf of client machines.
- Parallel databases
 - Many core shared memory
 - Shared disk
 - Shared nothing
- Distributed databases
 - Geographical distribution
 - Schema/data heterogeneity

Database Architecture (Centralized/Shared-Memory)




Database Applications

Database applications are usually partitioned into two or

- Two-tier architecture -- the application resides at the where it invokes database system functionality at the
- Three-tier architecture -- the client machine acts as a does not contain any direct database calls.
 - The client end communicates with an application through a forms interface.
 - The application server in turn communicates with system to access data.


Two-tier and three-tier architecture



(a) Two-tier architecture

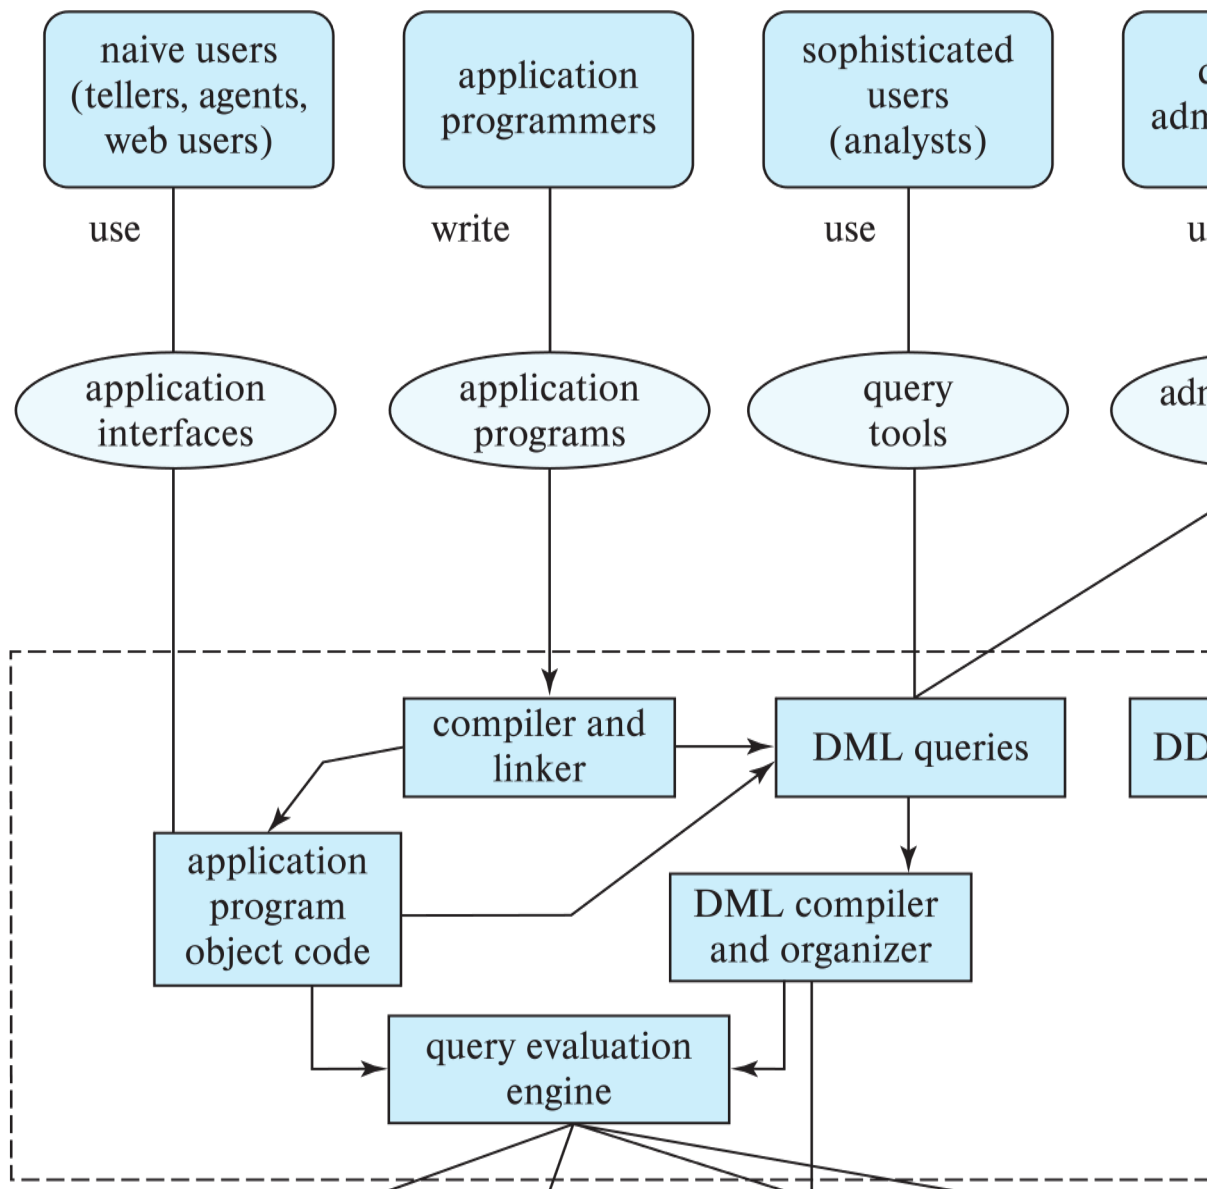
client

server



(b) Three-tier architecture

Database Users May See the Data D



Database Administrator

A person who has central control over the system is called **administrator (DBA)**. Functions of a DBA include:

- Schema definition
- Storage structure and access-method definition
- Schema and physical-organization modification
- Granting of authorization for data access
- Routine maintenance
- Periodically backing up the database
- Ensuring that enough free disk space is available for operations, and upgrading disk space as required
- Monitoring jobs running on the database

History of Database Systems

- 1950s and early 1960s:
 - Data processing using magnetic tapes for storage
 - Tapes provided only sequential access
 - Punched cards for input
- Late 1960s and 1970s:
 - Hard disks allowed direct access to data
 - Network and hierarchical data models in widespread use
 - Ted Codd defines the relational data model
 - Would win the ACM Turing Award for this work
 - IBM Research begins System R prototype
 - UC Berkeley (Michael Stonebraker) begins Ingres
 - Oracle releases first commercial relational database
 - High-performance (for the era) transaction processing

History of Database Systems (Co

- 1980s:
 - Research relational prototypes evolve into commercial systems
 - SQL becomes industrial standard
 - Parallel and distributed database systems
 - Wisconsin, IBM, Teradata
 - Object-oriented database systems
- 1990s:
 - Large decision support and data-mining applications
 - Large multi-terabyte data warehouses
 - Emergence of Web commerce

History of Database Systems (Co

- 2000s
 - Big data storage systems
 - Google BigTable, Yahoo PNuts, Amazon,
 - “NoSQL” systems.
 - Big data analysis: beyond SQL
 - Map reduce and friends
- 2010s
 - SQL reloaded
 - SQL front end to Map Reduce systems
 - Massively parallel database systems
 - Multi-core main-memory databases